

PixelCNN Autoregressive Model

```
In [1]: %load_ext autoreload
%autoreload 2

import warnings
warnings.filterwarnings("ignore")

import tensorflow as tf
from tensorflow.keras import (layers, models, callbacks, optimizers, datasets)

import numpy as np
import matplotlib.pyplot as plt
from IPython.display import DisplayObject, display, Image
```

```
2024-04-07 15:34:23.243349: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.
2024-04-07 15:34:23.879816: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT
```

```
In [2]: IMAGE_SIZE=28
PIXEL_LEVELS = 4
N_FILTERS = 128
RESIDUAL_BLOCKS = 5
BATCH_SIZE = 128
EPOCHS = 25
```

```
In [3]: (x_train, _), (_, _) = datasets.fashion_mnist.load_data()

def preprocess(imgs_int):
    imgs_int = np.expand_dims(imgs_int, -1)
    imgs_int = tf.image.resize(imgs_int, (IMAGE_SIZE, IMAGE_SIZE)).numpy()
    imgs_int = (imgs_int / (256 / PIXEL_LEVELS)).astype(int)
    imgs = imgs_int.astype("float32")
    imgs = imgs / PIXEL_LEVELS
    return imgs, imgs_int

input_data, output_data = preprocess(x_train)
display(Image(input_data))

#plt.imshow(input_data)
```

```
2024-04-07 15:34:25.641564: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.081439: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.081489: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.106972: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.107030: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.107057: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.251771: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.251827: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.251835: I tensorflow/core/common_runtime/gpu/gpu_devic
e.cc:2019] Could not identify NUMA node of platform GPU id 0, defaulting t
o 0. Your kernel may not have been built with NUMA support.
2024-04-07 15:34:26.251869: I external/local_xla/xla/stream_executor/cuda/
cuda_executor.cc:984] could not open file to read NUMA node: /sys/bus/pci/
devices/0000:01:00.0/numa_node
Your kernel may have been built without NUMA support.
2024-04-07 15:34:26.251889: I tensorflow/core/common_runtime/gpu/gpu_devic
e.cc:1928] Created device /job:localhost/replica:0/task:0/device:GPU:0 wit
h 2248 MB memory: -> device: 0, name: NVIDIA GeForce GTX 1650, pci bus id
: 0000:01:00.0, compute capability: 7.5
```

```

In [4]: class MaskedConv2D(layers.Layer):
    def __init__(self, mask_type, **kwargs):
        super(MaskedConv2D, self).__init__()
        self.mask_type = mask_type
        self.conv = layers.Conv2D(**kwargs)

    def build(self, input_shape):
        self.conv.build(input_shape)
        kernel_shape = self.conv.kernel.shape
        self.mask = np.zeros(shape=kernel_shape)
        self.mask[:, kernel_shape[0] // 2, ...] = 1.0
        self.mask[kernel_shape[0] // 2, : kernel_shape[1] // 2, ...] = 1.0
        if self.mask_type == "B":
            self.mask[kernel_shape[0] // 2, kernel_shape[1] // 2, ...] = 1.0

    def call(self, inputs):
        self.conv.kernel.assign(self.conv.kernel * self.mask)
        return self.conv(inputs)

    def get_config(self):
        cfg = super().get_config()
        return cfg

class ResidualBlock(layers.Layer):
    def __init__(self, filters, **kwargs):
        super(ResidualBlock, self).__init__(**kwargs)
        self.conv1 = layers.Conv2D(filters=filters // 2, kernel_size=1, activation='relu')
        self.pixel_conv = MaskedConv2D(mask_type = "B", filters = filters // 2, kernel_size=1, activation='relu')
        self.conv2 = layers.Conv2D(filters=filters, kernel_size=1, activation='relu')

    def call(self, inputs):
        x = self.conv1(inputs)
        x = self.pixel_conv(x)
        x = self.conv2(x)
        return layers.add([inputs, x])

    def get_config(self):
        cfg = super().get_config()
        return cfg

```

```
In [5]: inputs = layers.Input(shape=(IMAGE_SIZE, IMAGE_SIZE, 1))
x = MaskedConv2D(mask_type="A", filters=N_FILTERS, kernel_size=3, activation='relu')(inputs)

for _ in range(RESIDUAL_BLOCKS):
    x = ResidualBlock(filters=N_FILTERS)(x)

for _ in range(2):
    x = MaskedConv2D(mask_type="B", filters=N_FILTERS, kernel_size=1, stride=2)(x)

out = layers.Conv2D(filters = PIXEL_LEVELS, kernel_size=1, strides=1, activation='relu')(x)

pixel_cnn = models.Model(inputs, out)
pixel_cnn.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 28, 28, 1)	0
masked_conv2d (MaskedConv2D)	(None, 28, 28, 128)	1,280
residual_block (ResidualBlock)	(None, 28, 28, 128)	53,504
residual_block_1 (ResidualBlock)	(None, 28, 28, 128)	53,504
residual_block_2 (ResidualBlock)	(None, 28, 28, 128)	53,504
residual_block_3 (ResidualBlock)	(None, 28, 28, 128)	53,504
residual_block_4 (ResidualBlock)	(None, 28, 28, 128)	53,504
masked_conv2d_6 (MaskedConv2D)	(None, 28, 28, 128)	16,512
masked_conv2d_7 (MaskedConv2D)	(None, 28, 28, 128)	16,512
conv2d_18 (Conv2D)	(None, 28, 28, 4)	516

Total params: 302,340 (1.15 MB)

Trainable params: 302,340 (1.15 MB)

Non-trainable params: 0 (0.00 B)

```
In [6]: adam = optimizers.Adam(learning_rate=0.0005)
loss = losses.SparseCategoricalCrossentropy(tf.keras.losses.Reduction.SUM)
pixel_cnn.compile(loss=loss, optimizer=adam, metrics=["accuracy", losses.SparseCategoricalAccuracy])
```

```

In [7]: tensorboard_callback = callbacks.TensorBoard(log_dir = "./logs")
        model_checkpoint_callback = callbacks.ModelCheckpoint(filepath="./checkpoint",
                                                                save_weights_only=True)

class ImageGenerator(callbacks.Callback):
    def __init__(self, num_img):
        self.num_img = num_img

    def sample_from(self, probs, temperature):
        probs = probs ** (1 / temperature)
        probs = probs / np.sum(probs)
        return np.random.choice(len(probs), p=probs)

    def generate(self, temperature):
        generated_images = np.zeros(shape=(self.num_img,) + (pixel_cnn.input_shape[1:],
                                                                batch, rows, cols, channels = generated_images.shape)
        for row in range(rows):
            for col in range(cols):
                for channel in range(channels):
                    probs = self.model.predict(generated_images, verbose=0)
                    generated_images[:, row, col, channel] = [self.sample_from(probs)]
                    generated_images[:, row, col, channel] /= PIXEL_LEVELS
        return generated_images

    def on_epoch_end(self, epoch, logs=None):
        generated_images = self.generate(temperature=1.0)
        plt.figure(figsize=(5,5))
        for i in generated_images:
            plt.imshow(i)
            #plt.savefig(f"Generated_images_{i}", format="png")
image_generator_callback = ImageGenerator(num_img=10)

```

```
In [8]: # steps_per_epoch=1 to train faster but results are non accurate
pixel_cnn.fit(input_data, output_data, batch_size=BATCH_SIZE, epochs=EPOCHS)
```

Epoch 1/25

WARNING: All log messages before absl::InitializeLog() is called are written to STDERR

I0000 00:00:1712484273.469948 666377 service.cc:145] XLA service 0x7fa84400ef20 initialized for platform CUDA (this does not guarantee that XLA will be used). Devices:

I0000 00:00:1712484273.470016 666377 service.cc:153] StreamExecutor device (0): NVIDIA GeForce GTX 1650, Compute Capability 7.5


2024-04-07 15:34:33.568880: I tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:268] disabling MLIR crash reproducer, set env var `MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.

W0000 00:00:1712484273.750199 666377 assert_op.cc:38] Ignoring Assert operator compile_loss/sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

W0000 00:00:1712484273.977995 666377 assert_op.cc:38] Ignoring Assert operator sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

2024-04-07 15:34:34.231257: I external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:465] Loaded cuDNN version 8902

I0000 00:00:1712484281.064228 666377 device_compiler.h:188] Compiled cluster using XLA! This line is logged at most once for the lifetime of the process.


1/1  **0s** 11s/step - accuracy: 0.6129 - loss: 1.3869 - sparse_categorical_crossentropy: 1.3869

W0000 00:00:1712484281.875514 666374 assert_op.cc:38] Ignoring Assert operator compile_loss/sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert


W0000 00:00:1712484281.883013 666374 assert_op.cc:38] Ignoring Assert operator sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

W0000 00:00:1712484283.635918 666376 assert_op.cc:38] Ignoring Assert operator compile_loss/sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert


W0000 00:00:1712484283.639935 666376 assert_op.cc:38] Ignoring Assert operator sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

1/1  **56s** 56s/step - accuracy: 0.6129 - loss: 1.3869 - sparse_categorical_crossentropy: 1.3869 - val_accuracy: 0.7445 - val_loss: 1.3738 - val_sparse_categorical_crossentropy: 1.3738


Epoch 2/25

1/1  **41s** 41s/step - accuracy: 0.7470 - loss: 1.3736 - sparse_categorical_crossentropy: 1.3736 - val_accuracy: 0.7725 - val_loss: 1.3602 - val_sparse_categorical_crossentropy: 1.3602

Epoch 3/25

1/1  **43s** 43s/step - accuracy: 0.7702 - loss: 1.3604 - sparse_categorical_crossentropy: 1.3604 - val_accuracy: 0.7736 - val_loss: 1.3458 - val_sparse_categorical_crossentropy: 1.3458

Epoch 4/25

1/1  **43s** 43s/step - accuracy: 0.7761 - loss: 1.3464 - sparse_categorical_crossentropy: 1.3464 - val_accuracy: 0.7734 - val_loss: 1.3295 - val_sparse_categorical_crossentropy: 1.3295


Epoch 5/25



1/1  **43s** 43s/step - accuracy: 0.7627 - loss: 1.3295 -



sparse_categorical_crossentropy: 1.3295 - val_accuracy: 0.7733 - val_loss: 1.3112 - val_sparse_categorical_crossentropy: 1.3112
Epoch 6/25
1/1 ————— 43s 43s/step - accuracy: 0.7655 - loss: 1.3128 - sparse_categorical_crossentropy: 1.3128 - val_accuracy: 0.7733 - val_loss: 1.2906 - val_sparse_categorical_crossentropy: 1.2906
Epoch 7/25
1/1 ————— 42s 42s/step - accuracy: 0.7761 - loss: 1.2904 - sparse_categorical_crossentropy: 1.2904 - val_accuracy: 0.7733 - val_loss: 1.2675 - val_sparse_categorical_crossentropy: 1.2675
Epoch 8/25
1/1 ————— 42s 42s/step - accuracy: 0.7892 - loss: 1.2630 - sparse_categorical_crossentropy: 1.2630 - val_accuracy: 0.7731 - val_loss: 1.2407 - val_sparse_categorical_crossentropy: 1.2407
Epoch 9/25
1/1 ————— 43s 43s/step - accuracy: 0.7837 - loss: 1.2390 - sparse_categorical_crossentropy: 1.2390 - val_accuracy: 0.7731 - val_loss: 1.2100 - val_sparse_categorical_crossentropy: 1.2100
Epoch 10/25
1/1 ————— 42s 42s/step - accuracy: 0.7728 - loss: 1.2095 - sparse_categorical_crossentropy: 1.2095 - val_accuracy: 0.7731 - val_loss: 1.1747 - val_sparse_categorical_crossentropy: 1.1747
Epoch 11/25
1/1 ————— 43s 43s/step - accuracy: 0.7887 - loss: 1.1655 - sparse_categorical_crossentropy: 1.1655 - val_accuracy: 0.7731 - val_loss: 1.1341 - val_sparse_categorical_crossentropy: 1.1341
Epoch 12/25
1/1 ————— 43s 43s/step - accuracy: 0.7890 - loss: 1.1246 - sparse_categorical_crossentropy: 1.1246 - val_accuracy: 0.7730 - val_loss: 1.0877 - val_sparse_categorical_crossentropy: 1.0877
Epoch 13/25
1/1 ————— 44s 44s/step - accuracy: 0.7606 - loss: 1.0951 - sparse_categorical_crossentropy: 1.0951 - val_accuracy: 0.7730 - val_loss: 1.0351 - val_sparse_categorical_crossentropy: 1.0351
Epoch 14/25
1/1 ————— 42s 42s/step - accuracy: 0.7585 - loss: 1.0492 - sparse_categorical_crossentropy: 1.0492 - val_accuracy: 0.7729 - val_loss: 0.9772 - val_sparse_categorical_crossentropy: 0.9772
Epoch 15/25
1/1 ————— 40s 40s/step - accuracy: 0.7863 - loss: 0.9616 - sparse_categorical_crossentropy: 0.9616 - val_accuracy: 0.7725 - val_loss: 0.9143 - val_sparse_categorical_crossentropy: 0.9143
Epoch 16/25
1/1 ————— 41s 41s/step - accuracy: 0.7767 - loss: 0.9076 - sparse_categorical_crossentropy: 0.9076 - val_accuracy: 0.7723 - val_loss: 0.8491 - val_sparse_categorical_crossentropy: 0.8491
Epoch 17/25
1/1 ————— 39s 39s/step - accuracy: 0.7843 - loss: 0.8365 - sparse_categorical_crossentropy: 0.8365 - val_accuracy: 0.7723 - val_loss: 0.7845 - val_sparse_categorical_crossentropy: 0.7845
Epoch 18/25
1/1 ————— 39s 39s/step - accuracy: 0.7664 - loss: 0.7854 - sparse_categorical_crossentropy: 0.7854 - val_accuracy: 0.7723 - val_loss: 0.7242 - val_sparse_categorical_crossentropy: 0.7242
Epoch 19/25
1/1 ————— 41s 41s/step - accuracy: 0.7629 - loss: 0.7440 - sparse_categorical_crossentropy: 0.7440 - val_accuracy: 0.7722 - val_loss: 0.6714 - val_sparse_categorical_crossentropy: 0.6714
Epoch 20/25
1/1 ————— 40s 40s/step - accuracy: 0.7817 - loss: 0.6465 -




```
sparse_categorical_crossentropy: 0.6465 - val_accuracy: 0.7721 - val_loss: 0.6279 - val_sparse_categorical_crossentropy: 0.6279
Epoch 21/25
1/1 ————— 46s 46s/step - accuracy: 0.7679 - loss: 0.6383 - sparse_categorical_crossentropy: 0.6383 - val_accuracy: 0.7720 - val_loss: 0.5936 - val_sparse_categorical_crossentropy: 0.5936
Epoch 22/25
1/1 ————— 41s 41s/step - accuracy: 0.7747 - loss: 0.6040 - sparse_categorical_crossentropy: 0.6040 - val_accuracy: 0.7716 - val_loss: 0.5651 - val_sparse_categorical_crossentropy: 0.5651
Epoch 23/25
1/1 ————— 44s 44s/step - accuracy: 0.7739 - loss: 0.5610 - sparse_categorical_crossentropy: 0.5610 - val_accuracy: 0.7711 - val_loss: 0.5411 - val_sparse_categorical_crossentropy: 0.5411
Epoch 24/25
1/1 ————— 41s 41s/step - accuracy: 0.7657 - loss: 0.5596 - sparse_categorical_crossentropy: 0.5596 - val_accuracy: 0.7699 - val_loss: 0.5206 - val_sparse_categorical_crossentropy: 0.5206
Epoch 25/25
1/1 ————— 42s 42s/step - accuracy: 0.7752 - loss: 0.5018 - sparse_categorical_crossentropy: 0.5018 - val_accuracy: 0.7699 - val_loss: 0.5031 - val_sparse_categorical_crossentropy: 0.5031
```



Out[8]: <keras.src.callbacks.history.History at 0x7fa93477fc90>











































```
In [9]: generated_images = image_generator_callback.generate(temperature=1.0)
display(Image(generated_images))
```



```
In [21]: try:
pixel_cnn.save("./models/pixelcnn.keras", overwrite=True)
pixel_cnn.save("./models/pixel_cnn.h5", overwrite=True)
except Exception as e:
    print(str(e)[-1])
    print("Model Name already exists")
```

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy . We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

Unable to synchronously create dataset (name already exists)
Model Name already exists

```
In [11]: display(Image(generated_images))
```



```
In [12]: pixel_cnn.history.history.keys()
```

```
Out[12]: dict_keys([])
```

```
In [20]: def deprocess_image(x):
x = np.stack((x, x, x), 2)
x *= 255.0
x = np.clip(x, 0, 255).astype("uint8")
return x

generated_images = image_generator_callback.generate(temperature=1.0)

for i, pic in enumerate(generated_images):
    tf.keras.utils.save_img("./outputs/generative_outputs/pixelcnn_outputs/%i.png", pic)

for i in range(0,9):
    display(Image(f"./outputs/generative_outputs/pixelcnn_outputs/generated_{i}.png"))
```




In [14]: `%timeit pixel_cnn.fit`

33.7 ns \pm 0.205 ns per loop (mean \pm std. dev. of 7 runs, 10,000,000 loops each)

In [15]: `pixel_cnn.evaluate(input_data, output_data)`

W0000 00:00:1712485422.559829 666377 assert_op.cc:38] Ignoring Assert operator compile_loss/sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

W0000 00:00:1712485422.565238 666377 assert_op.cc:38] Ignoring Assert operator sparse_categorical_crossentropy/SparseSoftmaxCrossEntropyWithLogits/assert_equal_1/Assert/Assert

1875/1875  **17s** 9ms/step - accuracy: 0.7729 - loss: 0.4974 - sparse_categorical_crossentropy: 0.4974

Out[15]: [0.4993739128112793, 0.7715013027191162, 0.4993739128112793]

In [16]: `pixel_cnn.get_metrics_result()`

Out[16]: {'accuracy': 0.7715013027191162,
 'loss': 0.4993739128112793,
 'sparse_categorical_crossentropy': 0.4993739128112793}

In [17]: `pixel_cnn.history.history.keys()`

Out[17]: dict_keys([])

In []: